

CISC 3320 HW Assignment – 3 (7 pts)

This assignment accentuates the importance of synchronization when there is concurrency. You are asked to manufacture a situation where data inconsistency results from multiple threads trying to write into a common piece of data. Then you write another version of the same program that solves the issue present in the previous version (the difference is quite small)

Requirements

1. Use either C/C++ or Java to write two programs: SyncDemo1 and SyncDemo2
2. In SyncDemo1
 - Populate an array or `ArrayList` with integers 1~1000.
 - Calculate the sum of all its elements – you do this by creating 5 threads and passing to each thread 1/5 of the data (i.e., one adds 1~200, another adds 201~400, and so on...)
 - In HW-2 you are recommended to use threads that implement `Callable` so they can return their result. In this exercise, you need threads that implement `Runnable` and add their subtotal into a grand total, which is not protected with synchronization.
 - Your main thread then prints the grand total, and it should show an incorrect result (the correct sum of $1+2+\dots+1000 = 500,500$)
3. In SyncDemo2
 - Fix the data inconsistency issue exhibited in SyncDemo1 by employing synchronization measures, so that the sum reported is correct.

Hints

With C/C++ the grand total would be a global variable. With Java, one can use a static variable of the main class. Such a variable can be accessed by a static inner class which will be your task class that implements `Runnable`.

To create data race, you may have to exaggerate the random sequencing of the updating process of the `grandTotal` variable. So instead of writing `grandTotal += subtotal`, you can use:

```
int temp = grandTotal + subTotal;
// wait a few milliseconds
grandTotal = temp;
```

With Java, you wouldn't be able to solve the problem just by adding a `synchronized` keyword in front of your thread's worker method definition, if each thread instance is independent. That's because this feature makes use of the thread object's intrinsic lock, and 5 threads will have 5 independent intrinsic locks – losing protection for any individual thread. There are several possible solutions. You can explicitly use a single `ReentrantLock`. You can also give each thread the access to a common object, and make the `synchronized` function a member of that shared object so they all use the same intrinsic lock (this is the case with the `BoundedBuffer` example in the slides). Alternatively, note that the `synchronized` keyword can be used to create a *synchronized block*:

```
synchronized (expr) { // expr must evaluate to an object reference
    statements;
}
```

In this case, the statements inside the block will acquire the lock on the object represented by `expr` (which could be any shared object), not the individual thread instance. In all, using `ReentrantLock` is probably the easiest and most straightforward approach.

What to turn in

- Please submit (via email) the following:
 - your source code in separate files
 - running result (you could use a screen capture)
- Please do not copy code, whether from classmates or from the Internet.