

# Chapter 7: Synchronization Examples

---



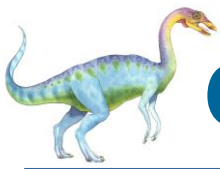


# Outline

---

- Explain the bounded-buffer synchronization problem
- Explain the readers-writers synchronization problem
- Explain and dining-philosophers synchronization problems
- Describe the tools used by Linux and Windows to solve synchronization problems.
- Illustrate how POSIX and Java can be used to solve process synchronization problems





# Classical Problems of Synchronization

---

- Classical problems used to test newly-proposed synchronization schemes
  - Bounded-Buffer Problem
  - Readers and Writers Problem
  - Dining-Philosophers Problem





# Bounded-Buffer Problem

---

- $n$  buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
- Semaphore **full** initialized to the value 0 (# of full buffers)
- Semaphore **empty** initialized to the value  $n$  (# of empty buffers)





# Bounded Buffer Problem (Cont.)

- The structure of the producer process

```
while (true) {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty) ;  
    wait(mutex) ;  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex) ;  
    signal(full) ;  
}
```

Protects against adding into full; empty will reach 0 (blocking) after n productions if consumer never consumes any, i.e., signal(empty)

lock for the CS

Protects against drawing from empty; full is incremented by 1 for every production; 0 represents no item to be consumed, causing the consumer process to wait





# Bounded Buffer Problem (Cont.)

- The structure of the consumer process

```
while (true) {  
    wait(full) ; Wait if # of full buffers is zero (nothing to consume)  
    wait(mutex) ; Lock the CS  
  
    ...  
    /* remove an item from buffer to next_consumed */  
    ...  
    signal(mutex) ;  
    signal(empty) ; Increase the # of empty buffers by 1  
  
    ...  
    /* consume the item in next consumed */  
    ...  
}
```





# Readers-Writers Problem

---

- A data set is shared among a number of concurrent processes
  - **Readers** – only read the data set; they do ***not*** perform any updates
  - **Writers** – can both read and write
- Problem – allow multiple readers to read at the same time
  - Only one single writer can access the shared data at the same time
- Several variations of how readers and writers are considered – all involve some form of priorities





# Readers-Writers Problem (Cont.)

---

- Shared Data
  - Data set
  - Semaphore **rw\_mutex** initialized to 1 Only one writer allowed at a time
  - Semaphore **mutex** initialized to 1
  - Integer **read\_count** initialized to 0

This is known as the *first readers-writers problem*: Access is favored to readers; writers can only access data while there are no readers, i.e., writers may starve.







# Readers-Writers Problem (Cont.)

- The structure of a writer process

```
while (true) {  
    wait(rw_mutex);  
  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
}
```





# Readers-Writers Problem (Cont.)

- The structure of a reader process

```
while (true){  
    wait(mutex); protects modifying read_count  
    read_count++;  
    if (read_count == 1) /* first reader */  
        wait(rw_mutex); only blocks writer access, more  
                           readers are always allowed  
    signal(mutex); opens access to read_count  
  
    ...  
    /* reading is performed */  
    ...  
    wait(mutex); again, protects modifying read_count  
    read_count--;  
    if (read_count == 0) /* last reader */  
        signal(rw_mutex); opens writer access if no readers  
    signal(mutex); restores access to read_count  
}
```





# Readers-Writers Problem Variations

---

- The solution in previous slide can result in a situation where a writer process never writes. It is referred to as the “First reader-writer” problem.
- The “Second reader-writer” problem is a variation of the first reader-writer problem that states:
  - Once a writer is ready to write, no “newly arrived reader” is allowed to read.
- Both the first and second may result in starvation, leading to even more variations
- Best suited to applications where
  - Readers and writers are easily identified
  - There are more readers than writers





# Kernel Synchronization - Windows


- Uses interrupt masks to protect access to global resources on uniprocessor systems
- Uses **spinlocks** on multiprocessor systems
  - For protecting short code segments
  - Kernel ensures a spinlocking-thread will never be preempted (spinlocks are generally used when the lock will be held for less than two context switches)
- For thread synchronization outside the kernel, Windows also provides user-land **dispatcher objects** which may act as mutexes, semaphores, events, and timers
  - **Events**
    - ▶ An event acts much like a condition variable
  - Timers notify one or more thread when time expired
  - Dispatcher objects may be in either **signaled-state** (object available) or **non-signaled state** (thread will block)





# Kernel Synchronization - Windows

- Mutex dispatcher object



- If a thread blocks on a nonsignaled dispatcher object, its state changes from ready to waiting
- When the state for the dispatcher object moves to signaled, the kernel can move **one thread** (waiting on a mutex), or **some threads** (waiting on semaphores), or **all threads** (waiting on events), from the waiting state to the ready state so they can resume executing





# Kernel Synchronization - Linux

---

- Linux:
  - Prior to kernel Version 2.6, disables interrupts to implement short critical sections (i.e., nonpreempted)
  - Version 2.6 and later, fully preemptive
- Linux provides:
  - Atomic integers (all math using them are not interrupted)
  - Mutex locks (`mutex_lock()` and `mutex_unlock()`)
  - Semaphores
  - Spinlocks
  - Reader-writer versions of both semaphores and spinlocks
- On SMP machines, spinlocks are used (other threads can still run on separate processors)
- On single-CPU system, spinlocks replaced by enabling and disabling kernel preemption





# Linux Synchronization

- Atomic variables

`atomic_t` is the type for atomic integer

- Consider the variables

```
atomic_t counter;  
int value;
```

<i>Atomic Operation</i>	<i>Effect</i>
<code>atomic_set(&amp;counter, 5);</code>	<code>counter = 5</code>
<code>atomic_add(10, &amp;counter);</code>	<code>counter = counter + 10</code>
<code>atomic_sub(4, &amp;counter);</code>	<code>counter = counter - 4</code>
<code>atomic_inc(&amp;counter);</code>	<code>counter = counter + 1</code>
<code>value = atomic_read(&amp;counter);</code>	<code>value = 12</code>

- Note that these methods pertain to synchronization available only to kernel developers. For user level programmers, use the POSIX API





# POSIX Synchronization

---

- POSIX API provides
  - mutex locks
  - semaphores
  - condition variable
- Widely used on UNIX, Linux, and macOS







# POSIX Mutex Locks

---

- Creating and initializing the lock

```
#include <pthread.h>
```

```
pthread_mutex_t mutex;
```

```
/* create and initialize the mutex lock */  
pthread_mutex_init(&mutex, NULL);
```

- Acquiring and releasing the lock

```
/* acquire the mutex lock */  
pthread_mutex_lock(&mutex);
```

```
/* critical section */
```

```
/* release the mutex lock */  
pthread_mutex_unlock(&mutex);
```





# POSIX Semaphores

---

- POSIX provides two versions – **named** and **unnamed**.
- Named semaphores can be used by unrelated processes, unnamed cannot.





# POSIX Named Semaphores

- Creating and initializing the semaphore:

```
#include <semaphore.h>
sem_t *sem;
```

```
/* Create the semaphore and initialize it to 1 */
sem = sem_open("SEM", O_CREAT, 0666, 1);
```

- Another process can access the semaphore by referring to its name **SEM**.
  - They can use `sem = sem_open("SEM", O_RDWR)`
- Acquiring and releasing the semaphore:

```
/* acquire the semaphore */
sem_wait(sem);
```

```
/* critical section */
```

```
/* release the semaphore */
sem_post(sem);
```





# POSIX Unnamed Semaphores

- Creating and initializing the semaphore:

```
#include <semaphore.h>
sem_t sem;  Related processes are those that have
            access to this global variable sem

/* Create the semaphore and initialize it to 1 */
sem_init(&sem, 0, 1); The second argument: 0: shared among
                    threads, 1: shared among processes
```

- Acquiring and releasing the semaphore:

```
/* acquire the semaphore */
sem_wait(&sem);

/* critical section */

/* release the semaphore */
sem_post(&sem);
```





# POSIX Condition Variables

---

- Since POSIX is typically used in C/C++ and these languages do not provide a monitor, POSIX condition variables are associated with a POSIX mutex lock to provide mutual exclusion: Creating and initializing the condition variable:

```
pthread_mutex_t mutex;  
pthread_cond_t cond_var;  
  
pthread_mutex_init(&mutex, NULL);  
pthread_cond_init(&cond_var, NULL);
```





# POSIX Condition Variables

- Thread waiting for the condition `a == b` to become true:

```
pthread_mutex_lock(&mutex);  
while (a != b)  
    pthread_cond_wait(&cond_var, &mutex);  
  
pthread_mutex_unlock(&mutex);
```

Here, it releases the mutex lock on entry, and blocks, then regains the mutex on successful return.

Can we use `if (a!=b)`?

- Thread signaling another thread waiting on the condition variable:

```
pthread_mutex_lock(&mutex);  
a = b;  
pthread_cond_signal(&cond_var);  
pthread_mutex_unlock(&mutex);
```





# Java Synchronization

---

- Java provides rich set of synchronization features:
  - Java monitors
  - Reentrant locks
  - Semaphores
  - Condition variables





# Java Monitors

---

- Every Java object has associated with it a single lock.
- If a method is declared as **synchronized**, a calling thread must own the lock for the object.
- If the lock is owned by another thread, the calling thread must wait for the lock until it is released.
- Locks are released when the owning thread exits the **synchronized** method.







# Bounded Buffer – Java Synchronization

```
public class BoundedBuffer<E>
{
    private static final int BUFFER_SIZE = 5;

    private int count, in, out;
    private E[] buffer;

    public BoundedBuffer() {
        count = 0;
        in = 0;
        out = 0;
        buffer = (E[]) new Object[BUFFER_SIZE];
    }

    /* Producers call this method */
    public synchronized void insert(E item) {
        /* See Figure 7.11 */ Slide 7.29
    }


    /* Consumers call this method */
    public synchronized E remove() {
        /* See Figure 7.11 */ Slide 7.30
    }
}
```





# Java Synchronization

- A thread that tries to acquire an unavailable lock is placed in the object's **entry set**:




When the lock is available, JVM arbitrarily picks one thread (usually FCFS though) to become the owner of the object lock and it can enter the *synchronized* method. The lock is released when the thread exits such method.





# Java Synchronization

- Similarly, each object also has a **wait set**.
- When a thread calls **wait()**:
  1. It releases the lock for the object
  2. The state of the thread is set to blocked
  3. The thread is placed in the wait set for the object



A thread inside the *synchronized* method (owning the lock) may have to wait for certain event (condition). For example, a producer calls `insert()` and the buffer is full. The thread will release the lock and wait inside the wait set until the condition (one buffer's content consumed) is met to continue.





# Java Synchronization

---

- A thread typically calls `wait()` when it is waiting for a condition to become true.
- How does a thread get notified?
- When a thread calls `notify()`:
  1. An arbitrary thread T is selected from the wait set
  2. T is moved from the wait set to the entry set
  3. Set the state of T from blocked to runnable.
- T can now compete for the lock to check if the condition it was waiting for is now true (so use while-loop instead of if-clause).





# Bounded Buffer – Java Synchronization

```
/* Producers call this method */
public synchronized void insert(E item) {
    while (count == BUFFER_SIZE) {
        try {
            lock is released, and calling thread (producer) blocks and
            being placed in the wait set; Will regain the lock and
            return when notified by consumer
            wait();
        }
        catch (InterruptedException ie) { }
    }

    buffer[in] = item;
    in = (in + 1) % BUFFER_SIZE;
    count++;

    notify();
}
```

Picks an arbitrary thread from the wait set, and moves it to the entry set, setting its state from *blocked* to *runnable*. If there is no thread in wait set, the call is ignored.





# Bounded Buffer – Java Synchronization

---

```
/* Consumers call this method */
public synchronized E remove() {
    E item;

    while (count == 0) {
        try {
            wait();
        }
        catch (InterruptedException ie) { }
    }

    item = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    count--;

    notify();

    return item;
}
```





# Java Reentrant Locks

- Similar to mutex locks (with a fairness feature that favors granting the lock to the longest-waiting thread – default is no particular order)
- The **finally** clause ensures the lock will be released in case an exception occurs in the **try** block.

```
Lock key = new ReentrantLock();
```

```
key.lock();  
try {  
    /* critical section */  
}  
finally {  
    key.unlock();  
}
```

Returns if the lock is available or already owned by the current thread – thus reentrant

In the **finally** clause to guarantee proper releasing of the lock in case of exceptions





# Java Semaphores

---

- Constructor:

```
Semaphore(int value);
```

- Usage:

```
Semaphore sem = new Semaphore(1);

try {
    sem.acquire();
    /* critical section */
}
catch (InterruptedException ie) { }
finally {
    sem.release();
}
```







# Java Condition Variables

---

- Condition variables are associated with an **ReentrantLock**.
- Creating a condition variable using **newCondition()** method of **ReentrantLock**:

```
Lock key = new ReentrantLock();  
Condition condVar = key.newCondition();
```

- A thread waits by calling the **await()** method, and signals by calling the **signal()** method.





# Java Condition Variables

- Example (a bunch of workers take turns to do intermittent work):
- Five threads numbered 0 .. 4
- Shared variable `turn` indicating which thread's turn it is.
- Thread calls `doWork()` when it wishes to do some work. (But it may only do work if it is their turn.)
- If not their turn, wait
- If their turn, do some work for awhile .....
- When completed, notify the thread whose turn is next.
- Necessary data structures:

```
Lock lock = new ReentrantLock();  
Condition[] condVars = new Condition[5];  
  
for (int i = 0; i < 5; i++)  
    condVars[i] = lock.newCondition();
```

Here the condition is  
the circling turn





# Java Condition Variables

```
/* threadNumber is the thread that wishes to do some work */
public void doWork(int threadNumber)
{
    lock.lock();

    try {
        /**
         * If it's not my turn, then wait
         * until I'm signaled.
         */
        if (threadNumber != turn)
            condVars[threadNumber].await();


        /**
         * Do some work for awhile ...
         */

        /**
         * Now signal to the next thread.
         */
        turn = (turn + 1) % 5;
        condVars[turn].signal();
    }
    catch (InterruptedException ie) { }
    finally {
        lock.unlock();
    }
}
```



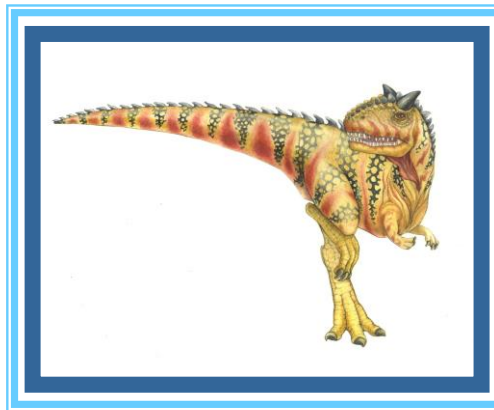
# End of Chapter 7

---



# Chapter 8: Deadlocks

---





# Outline

---

- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock





# Chapter Objectives

---

- Illustrate how deadlock can occur when mutex locks are used
- Define the four necessary conditions that characterize deadlock
- Identify a deadlock situation in a resource allocation graph
- Evaluate the four different approaches for preventing deadlocks
- Apply the banker's algorithm for deadlock avoidance
- Apply the deadlock detection algorithm
- Evaluate approaches for recovering from deadlock





# System Model

---

- System consists of resources
- Resource types  $R_1, R_2, \dots, R_m$ 
  - *CPU cycles, memory space, I/O devices*
- Each resource type  $R_i$  has  $W_i$  instances.
- Each process utilizes a resource as follows:
  - **request**
  - **use**
  - **release**







# Deadlock with Semaphores

---

- Data:
  - A semaphore **s1** initialized to 1
  - A semaphore **s2** initialized to 1
- Two processes **P1** and **P2**
- **P1**:  
`wait(s1)`  
`wait(s2)`
- **P2**:  
`wait(s2)`  
`wait(s1)`





# Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously.

- **Mutual exclusion:** only one process at a time can use a resource
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes
- **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task
- **Circular wait:** there exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .





# Resource-Allocation Graph

A set of vertices  $V$  and a set of edges  $E$ .


- $V$  is partitioned into two types:
  - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system
  - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system
- **request edge** – directed edge  $P_i \rightarrow R_j$
- **assignment edge** – directed edge  $R_j \rightarrow P_i$






# Resource Allocation Graph Example

- One instance of R1
- Two instances of R2
- One instance of R3
- Three instance of R4
- T1 holds one instance of R2 and is waiting for an instance of R1
- T2 holds one instance of R1, one instance of R2, and is waiting for an instance of R3
- T3 is holds one instance of R3






# Resource Allocation Graph with a Deadlock





# Graph with a Cycle But no Deadlock





# Basic Facts

---

- If graph contains no cycles  $\Rightarrow$  no deadlock
- If graph contains a cycle  $\Rightarrow$ 
  - if only one instance per resource type, then deadlock
  - if several instances per resource type, possibility of deadlock





# Methods for Handling Deadlocks

---

- Ensure that the system will **never** enter a deadlock state:
  - Deadlock prevention
  - Deadlock avoidance (avoid getting into unsafe states)
- Allow the system to enter a deadlock state and then recover
- Ignore the problem and pretend that deadlocks never occur in the system (this is the one used by most operating systems, including Linux and Windows)







# Deadlock Prevention

---

Invalidate one of the four necessary conditions for deadlock:

- **Mutual Exclusion** – not required for sharable resources (e.g., read-only files); must hold for non-sharable resources
- **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
  - Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none allocated to it.
  - Issues: Low resource utilization; starvation possible





# Deadlock Prevention (Cont.)

---

## ■ No Preemption:

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released
- Preempted resources are added to the list of resources for which the process is waiting
- Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting

## ■ Circular Wait:

- Impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration





# Circular Wait

- Invalidating the circular wait condition is most common.
- Simply assign each resource (i.e., mutex locks) a unique number.
- Resources must be acquired in order.
- If:

**F(first\_mutex) = 1**

**F(second\_mutex) = 5**

code for **thread\_two** could not be written as follows:



```
/* thread.one runs in this function */
void *do_work_one(void *param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);

    pthread_exit(0);
}
```

```
/* thread.two runs in this function */
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);

    pthread_exit(0);
}
```



# End of Chapter 8

---

